

# cm009 Exercises: tidy data

```
install.packages("dplyr")
```

```
library(tidyverse)
library(dplyr)
lotr <- read_csv("https://raw.githubusercontent.com/jennybc/lotr-tidy/master/data/lotr_tidy.csv")
guest <- read_csv("https://raw.githubusercontent.com/STAT545-UBC/Classroom/master/data/wedding/attend.csv")
email <- read_csv("https://raw.githubusercontent.com/STAT545-UBC/Classroom/master/data/wedding/emails.csv")
```

## Exercise 1: Univariate Pivoting

Consider the Lord of the Rings data:

```
lotr

## # A tibble: 18 x 4
##   Film                                Race  Gender Words
##   <chr>                                <chr> <chr>  <dbl>
## 1 The Fellowship Of The Ring Elf      Female 1229
## 2 The Fellowship Of The Ring Hobbit   Female  14
## 3 The Fellowship Of The Ring Man       Female   0
## 4 The Two Towers Elf              Female 331
## 5 The Two Towers Hobbit           Female   0
## 6 The Two Towers Man              Female 401
## 7 The Return Of The King Elf        Female 183
## 8 The Return Of The King Hobbit     Female   2
## 9 The Return Of The King Man         Female 268
## 10 The Fellowship Of The Ring Elf      Male   971
## 11 The Fellowship Of The Ring Hobbit   Male 3644
## 12 The Fellowship Of The Ring Man      Male 1995
## 13 The Two Towers Elf              Male   513
## 14 The Two Towers Hobbit           Male 2463
## 15 The Two Towers Man              Male 3589
## 16 The Return Of The King Elf        Male   510
## 17 The Return Of The King Hobbit     Male 2673
## 18 The Return Of The King Man         Male 2459
```

1. Would you say this data is in tidy format?
2. Widen the data so that we see the words spoken by each race, by putting race as its own column.

```
(lotr_wide <- lotr %>%
  pivot_wider(id_cols = c(-Race, -Words),
    names_from = Race,
    values_from = Words))
```

```
## # A tibble: 6 x 5
##   Film                                Gender  Elf Hobbit  Man
##   <chr>                                <chr>  <dbl> <dbl> <dbl>
## 1 The Fellowship Of The Ring Female 1229    14    0
## 2 The Two Towers Female 331      0 401
## 3 The Return Of The King Female 183      2 268
## 4 The Fellowship Of The Ring Male 971 3644 1995
```

```
## 5 The Two Towers           Male      513   2463  3589
## 6 The Return Of The King   Male      510   2673  2459
```

3. Re-lengthen the wide LOTR data from Question 2 above.

```
lotr_wide %>%
  pivot_longer(cols = c(-Film, -Gender),
               names_to = "Race",
               values_to = "Words")
```

```
## # A tibble: 18 x 4
##   Film                                Gender Race    Words
##   <chr>                                <chr> <chr> <dbl>
## 1 The Fellowship Of The Ring Female Elf     1229
## 2 The Fellowship Of The Ring Female Hobbit    14
## 3 The Fellowship Of The Ring Female Man         0
## 4 The Two Towers                Female Elf     331
## 5 The Two Towers                Female Hobbit     0
## 6 The Two Towers                Female Man     401
## 7 The Return Of The King        Female Elf     183
## 8 The Return Of The King        Female Hobbit     2
## 9 The Return Of The King        Female Man     268
## 10 The Fellowship Of The Ring Male    Elf     971
## 11 The Fellowship Of The Ring Male    Hobbit  3644
## 12 The Fellowship Of The Ring Male    Man    1995
## 13 The Two Towers                Male    Elf     513
## 14 The Two Towers                Male    Hobbit  2463
## 15 The Two Towers                Male    Man    3589
## 16 The Return Of The King        Male    Elf     510
## 17 The Return Of The King        Male    Hobbit  2673
## 18 The Return Of The King        Male    Man    2459
```

## Exercise 2: Multivariate Pivoting

Congratulations, you're getting married! In addition to the wedding, you've decided to hold two other events: a day-of brunch and a day-before round of golf. You've made a guestlist of attendance so far, along with food preference for the food events (wedding and brunch).

```
guest %>%
  DT::datatable(rownames = FALSE)
```

Show  entries Search:

party	name	meal_wedding	meal_brunch	attendance_wedding	attendance_brunch	attendance_golf
1	Sommer Medrano	PENDING	PENDING	PENDING	PENDING	PENDING
1	Phillip Medrano	vegetarian	Menu C	CONFIRMED	CONFIRMED	CONFIRMED
1	Blanka Medrano	chicken	Menu A	CONFIRMED	CONFIRMED	CONFIRMED
1	Emaan Medrano	PENDING	PENDING	PENDING	PENDING	PENDING
2	Blair Park	chicken	Menu C	CONFIRMED	CONFIRMED	CONFIRMED
2	Nigel Webb			CANCELLED	CANCELLED	CANCELLED
3	Sinead English	PENDING	PENDING	PENDING	PENDING	PENDING
4	Ayra Marks	vegetarian	Menu B	PENDING	PENDING	PENDING
5	Atlanta Connolly	PENDING	PENDING	PENDING	PENDING	PENDING
5	Denzel Connolly	fish	Menu B	CONFIRMED	CONFIRMED	CONFIRMED

Showing 1 to 10 of 30 entries Previous  2 3 Next

1. Put “meal” and “attendance” as their own columns, with the events living in a new column.

```
(guest_long <- guest %>%
  pivot_longer(cols = c(-party, -name),
    names_to = c(".value", "event"),
    names_sep = "_"))
```

```
## # A tibble: 90 x 5
##   party name          event meal attendance
##   <dbl> <chr>          <chr> <chr>    <chr>
## 1     1 Sommer Medrano wedding PENDING PENDING
## 2     1 Sommer Medrano brunch PENDING PENDING
## 3     1 Sommer Medrano golf <NA> PENDING
## 4     1 Phillip Medrano wedding vegetarian CONFIRMED
## 5     1 Phillip Medrano brunch Menu C CONFIRMED
## 6     1 Phillip Medrano golf <NA> CONFIRMED
## 7     1 Blanka Medrano wedding chicken CONFIRMED
## 8     1 Blanka Medrano brunch Menu A CONFIRMED
## 9     1 Blanka Medrano golf <NA> CONFIRMED
## 10    1 Emaan Medrano wedding PENDING PENDING
## # ... with 80 more rows
```

2. Use `tidyr::separate()` to split the name into two columns: “first” and “last”. Then, re-unite them with `tidyr::unite()`.

```
guest_long %>%
  separate(name, into = c("first", "last"), sep = " ") %>%
  unite(col = "name", first, last, sep = " ")
```

```
## # A tibble: 90 x 5
##   party name          event meal attendance
##   <dbl> <chr>          <chr> <chr>    <chr>
```

```
## 1      1 Sommer Medrano wedding PENDING PENDING
## 2      1 Sommer Medrano brunch PENDING PENDING
## 3      1 Sommer Medrano golf <NA> PENDING
## 4      1 Phillip Medrano wedding vegetarian CONFIRMED
## 5      1 Phillip Medrano brunch Menu C CONFIRMED
## 6      1 Phillip Medrano golf <NA> CONFIRMED
## 7      1 Blanka Medrano wedding chicken CONFIRMED
## 8      1 Blanka Medrano brunch Menu A CONFIRMED
## 9      1 Blanka Medrano golf <NA> CONFIRMED
## 10     1 Emaan Medrano wedding PENDING PENDING
## # ... with 80 more rows
```

3. Which parties still have a “PENDING” status for all members and all events?

```
guest_long %>%
  group_by(party) %>%
  summarize(all_pending = all(attendance == "PENDING"))
```

```
## # A tibble: 15 x 2
##   party all_pending
##   <dbl> <lgl>
## 1      1 FALSE
## 2      2 FALSE
## 3      3 TRUE
## 4      4 TRUE
## 5      5 FALSE
## 6      6 FALSE
## 7      7 FALSE
## 8      8 TRUE
## 9      9 FALSE
## 10     10 TRUE
## 11     11 FALSE
## 12     12 FALSE
## 13     13 FALSE
## 14     14 FALSE
## 15     15 FALSE
```

4. Which parties still have a “PENDING” status for all members for the wedding?

```
guest %>%
  group_by(party) %>%
  summarize(pending_wedding = all(attendance_wedding == "PENDING"))
```

```
## # A tibble: 15 x 2
##   party pending_wedding
##   <dbl> <lgl>
## 1      1 FALSE
## 2      2 FALSE
## 3      3 TRUE
## 4      4 TRUE
## 5      5 FALSE
## 6      6 FALSE
## 7      7 FALSE
## 8      8 TRUE
## 9      9 FALSE
## 10     10 TRUE
```

```
## 11    11 FALSE
## 12    12 FALSE
## 13    13 FALSE
## 14    14 FALSE
## 15    15 FALSE
```

5. Put the data back to the way it was.

```
guest_long %>%
  pivot_wider(id_cols = c(party, name),
              names_from = c(event),
              names_sep = "_",
              values_from = c(meal, attendance))

## # A tibble: 30 x 8
##   party name meal_wedding meal_brunch meal_golf attendance_wedd~
##   <dbl> <chr> <chr>         <chr>         <chr>         <chr>
## 1     1 Somme~ PENDING      PENDING      <NA>         PENDING
## 2     1 Phill~ vegetarian Menu C        <NA>         CONFIRMED
## 3     1 Blanka~ chicken    Menu A        <NA>         CONFIRMED
## 4     1 Emaan~ PENDING      PENDING      <NA>         PENDING
## 5     2 Blair~ chicken    Menu C        <NA>         CONFIRMED
## 6     2 Nigel~ <NA>         <NA>         <NA>         CANCELLED
## 7     3 Sinead~ PENDING      PENDING      <NA>         PENDING
## 8     4 Ayra~ vegetarian Menu B        <NA>         PENDING
## 9     5 Atla~ PENDING      PENDING      <NA>         PENDING
## 10    5 Denz~ fish        Menu B        <NA>         CONFIRMED
## # ... with 20 more rows, and 2 more variables: attendance_brunch <chr>,
## #   attendance_golf <chr>
```

6. You also have a list of emails for each party, in this worksheet under the variable `email`. Change this so that each person gets their own row. Use `tidyr::separate_rows()`

```
email %>%
  separate_rows(guest, sep = ", ")

## # A tibble: 28 x 2
##   guest      email
##   <chr>      <chr>
## 1 Sommer Medrano sommm@gmail.com
## 2 Phillip Medrano sommm@gmail.com
## 3 Blanka Medrano sommm@gmail.com
## 4 Emaan Medrano sommm@gmail.com
## 5 Blair Park bpark@gmail.com
## 6 Nigel Webb bpark@gmail.com
## 7 Sinead English singlish@hotmail.ca
## 8 Ayra Marks marksa42@gmail.com
## 9 Jolene Welsh jw1987@hotmail.com
## 10 Hayley Booker jw1987@hotmail.com
## # ... with 18 more rows
```

### Exercise 3: Making tibbles

1. Create a tibble that has the following columns:
  - A label column with "Sample A" in its entries.

- 100 random observations drawn from the  $N(0,1)$  distribution in the column `x`
- `y` calculated as the `x` values +  $N(0,1)$  error.

```
n <- 100
tibble(label = "Sample A",
        x = rnorm(n),
        y = x + rnorm(n))
```

```
## # A tibble: 100 x 3
##   label      x      y
##   <chr>    <dbl>  <dbl>
## 1 Sample A  1.02   1.88
## 2 Sample A -0.641  0.320
## 3 Sample A -0.268  1.16
## 4 Sample A -0.462 -0.803
## 5 Sample A  0.755  1.07
## 6 Sample A  0.218  1.24
## 7 Sample A -1.53  -0.00285
## 8 Sample A  0.221  1.69
## 9 Sample A -0.288  0.166
## 10 Sample A -0.346  0.0192
## # ... with 90 more rows
```

2. Generate a Gaussian sample of size 100 for each combination of the following means (`mu`) and standard deviations (`sd`).

```
n <- 100
mu <- c(-5, 0, 5)
sd <- c(1, 3, 10)
tibble(mu = mu, sd = sd) %>%
  group_by_all() %>%
  mutate(z = list(rnorm(n, mu, sd))) %>%
  unnest(z)
```

```
## # A tibble: 300 x 3
## # Groups:   mu, sd [3]
##      mu    sd    z
##    <dbl> <dbl> <dbl>
## 1    -5     1 -3.31
## 2    -5     1 -6.63
## 3    -5     1 -4.32
## 4    -5     1 -4.61
## 5    -5     1 -4.45
## 6    -5     1 -3.99
## 7    -5     1 -3.17
## 8    -5     1 -4.65
## 9    -5     1 -3.74
## 10   -5     1 -5.11
## # ... with 290 more rows
```

3. Fix the experiment tibble below (originally defined in the documentation of the `tidyr::expand()` function) so that all three repeats are displayed for each person, and the measurements are kept. The code is given, but needs one adjustment. What is it?

```
experiment <- tibble(
  name = rep(c("Alex", "Robert", "Sam"), c(3, 2, 1)),
  trt = rep(c("a", "b", "a"), c(3, 2, 1)),
```

```

rep = c(1, 2, 3, 1, 2, 1),
measurement_1 = runif(6),
measurement_2 = runif(6)
)
experiment %>% complete(nesting(name, trt), rep)

```

```

## # A tibble: 9 x 5
##   name    trt      rep measurement_1 measurement_2
##   <chr> <chr> <dbl>         <dbl>         <dbl>
## 1 Alex   a       1         0.439         0.372
## 2 Alex   a       2         0.541         0.340
## 3 Alex   a       3         0.285         0.836
## 4 Robert b       1         0.648         0.603
## 5 Robert b       2         0.309         0.712
## 6 Robert b       3          NA          NA
## 7 Sam    a       1         0.171         0.602
## 8 Sam    a       2          NA          NA
## 9 Sam    a       3          NA          NA

```